

## Sorting Strategies in Spark (HTX xData)

### Context

Large camera streaming datasets require efficient joins and ordering to support enrichment and alert ranking use-cases.

Dataset A (events) — millions per hour

Dataset B (locations) — ~10 K static records

### Sorting Strategies

#### Sort Within Partition

- `df.sortWithinPartitions("ts")`
- Cheap; preserves locality; used for window ops inside micro-batch.

#### Global Sort (`orderBy`)

- Full shuffle.
- Only when absolutely needed.
- Often inefficient for streaming.

#### Range Partitioning + Sort

- `repartitionByRange("location").sortWithinPartitions(...)`
- Best for range window analytics or distributed ordering.

#### Sort-Merge Join (SMJ)

- Required when join keys are large and broadcast is impossible.
- Must sort both sides + shuffle — expensive but scalable.

#### Broadcast Hash Join (BHJ)

- Ideal for small dimension table (~10-100 k rows).
- Zero shuffle join path.
- Default choice for this case.

### Strategy for This Project

- Broadcast dataset B to avoid shuffling camera events.
- If broadcast disabled (e.g., size growth), use SMJ on geographical\_location\_oid.

#### Additional Notes

- Use adaptive query execution (AQE) to auto-switch join strategy.
- Enable spark.sql.autoBroadcastJoinThreshold.
- Monitor skew; if hotspots appear, apply salting on hot OIDs.